Hugging Face

AMD DEVELOPER HACKATHON · MAY 11, 2026

REPOMIND

Open-source repo-scale coding agent on a single AMD MI300X.

256K context · FP8 · MIT · self-hosted · production-verified

Sardor Razikov · Tashkent

Independent ML Engineer · Top 1.9% Kaggle SPR · Top 22% AIMO3 · Zenodo author

Track 1 (AI Agents) · Hugging Face Special Prize · Build-in-Public · Qwen partner

`github.com/SRKRZ23/repomind · vLLM 0.17.1 + ROCm 7.2 · Qwen3-Coder-Next-FP8`

THE WAVE BIG TECH MISSED · $30B/YR LOCKED MARKET · ZERO INCUMBENTS

AMD

3 million developers locked out of AI coding tools.

Cursor ($9B). Anthropic ($40B). OpenAI ($500B). All blocked.

These headlines are public. Your CEO has seen them. The pain is real and unresolved 3 years later:

JPMorgan restricts employee use of ChatGPT

CNN · Feb 22, 2023
JP Morgan: 50,000 devs · banned ChatGPT · still in force

macrumors.com/2023/05/19/apple-bans-employees-from-using-chatgpt/

MacRumors / WSJ · May 19, 2023
Apple: 10K iOS devs · banned ChatGPT + Copilot

DEFENSESCOOP

DOD wants AI-enabled coding tools for 'tens of thousands' of users in its developer workforce

The products would enable AI-driven code generation, optimization, debugging, support and refinement at the edge.

DefenseScoop · Feb 26, 2026
DoD: "tens of thousands" devs · on-prem only · Feb 2026

THE LOCKED MARKET — VERIFIED HEADCOUNT:

JPM 50K devs · Apple 10K iOS · DoD tens of thousands · Pharma R&D · Sovereign-AI funds

$30B/yr productivity unlock that $9B Cursor + $40B Anthropic legally cannot reach.

REPOMIND is the open-source MIT solution for teams the SaaS market cannot serve.

Sources: cnn.com/2023/02/22 · macrumors.com/2023/05/19 · defensescoop.com/2026/02/26 · all primary, all verifiable, all linked from the public repo

AMD'S OWN BLOG SAID THIS EXISTS · WE SHIPPED THE OPEN-SOURCE PROOF

AMD

AMD set the trap.
REPOMIND walked through.

Run Qwen3-Coder-Next with vLLM on AMD Instinct GPUs

The integration of ROCm™ 7 software and vLLM allows users to fully exploit the 192GB HBM capacity of the MI300X GPU.

- **Unmatched Context:** Users can serve the full 256k context length on a single GPU using FP8 precision, a critical requirement for repo-level coding tasks that often exceed the memory limits of lesser hardware.
- **Optimized Throughput:** By leveraging tensor parallelism, developers can achieve the low-latency response times required for real-time IDE integrations like Claude Code or Trae.

AMD official · Feb 4, 2026 · amd.com/en/developer/resources/technical-articles/2026/day-0-support-for-qwen3-coder-next

VERIFIED 2026-05-05/06:

~143 GB
total workload (weights + KV + activations)

192 GB
MI300X single-card capacity

80 GB
H100 — short by 63 GB

2.4×
more memory single-card

92%
peak VRAM (176/191.7 GiB)

AMD'S OWN BLOG — WHAT WE TURNED INTO PROOF:

256K context · single GPU · FP8 precision — measured end-to-end on a real MI300X.

Mahdi Ghodsi (AMD AI Solution Architect) co-authored the blog · we filed AITER regression upstream · ROCm team gets reproducible bug report

VERIFIED ON REAL HARDWARE · 2026-05-05 / 06

It works. Here are the numbers.

Single MI300X x1 · vLLM 0.17.1 + ROCm 7.2 · 124 min total · 2 sessions · $4.12 cost

77.29 GiB
model weights in VRAM

94.58 GiB
KV cache available (2,065,744 tokens)

92%
VRAM peak (176 / 191.7 GiB)

31.31×
vLLM max concurrency at 256K

--max-model-len 262144 startup	**Application startup complete ✓**
/v1/models endpoint response	`max_model_len: 262144`
vLLM startup time (verified vllm_default.log:92)	**42.2 sec (vs industry months)**
Warm restart (model cached)	**1 min 22 sec**
Tuning attempt (AITER backend)	measured A/B regression — see slide 7

FULL EVIDENCE PACK (PUBLIC, MIT):

7 JSON results · 5 PNG plots · 15 e2e prompt/answer files · 2× rocm-smi snapshots · 3× vLLM startup logs

`github.com/SRKRZ23/repomind/tree/main/benchmarks/2026-05-05-mi300x-stress-test`

TIME-TO-WORKLOAD · LISA SU's "MONTHS → DAYS" KPI EXCEEDED

AMD

From git clone to working endpoint: 3 min 30 sec.

LISA SU · a16z FIRESIDE · OCT 2024:

"We need to compress optimization cycles from months down to days."

TIME TO PRODUCTION-READY 256K-CONTEXT INFERENCE:

42.2-SEC BREAKDOWN (vllm_default.log:75-92):

`model_load 28.45s + torch.compile 20.06s + graph_capture 21s → Application startup complete ✓`

Source: vllm_default.log — verified 2026-05-06 via "tail /tmp/vllm_default.log" → "INFO core.py:282 init engine took 42.20 seconds"

THROUGHPUT · 6 CONTEXTS · HOT MEASUREMENTS

TTFT linear in prompt size — prefill-bound.

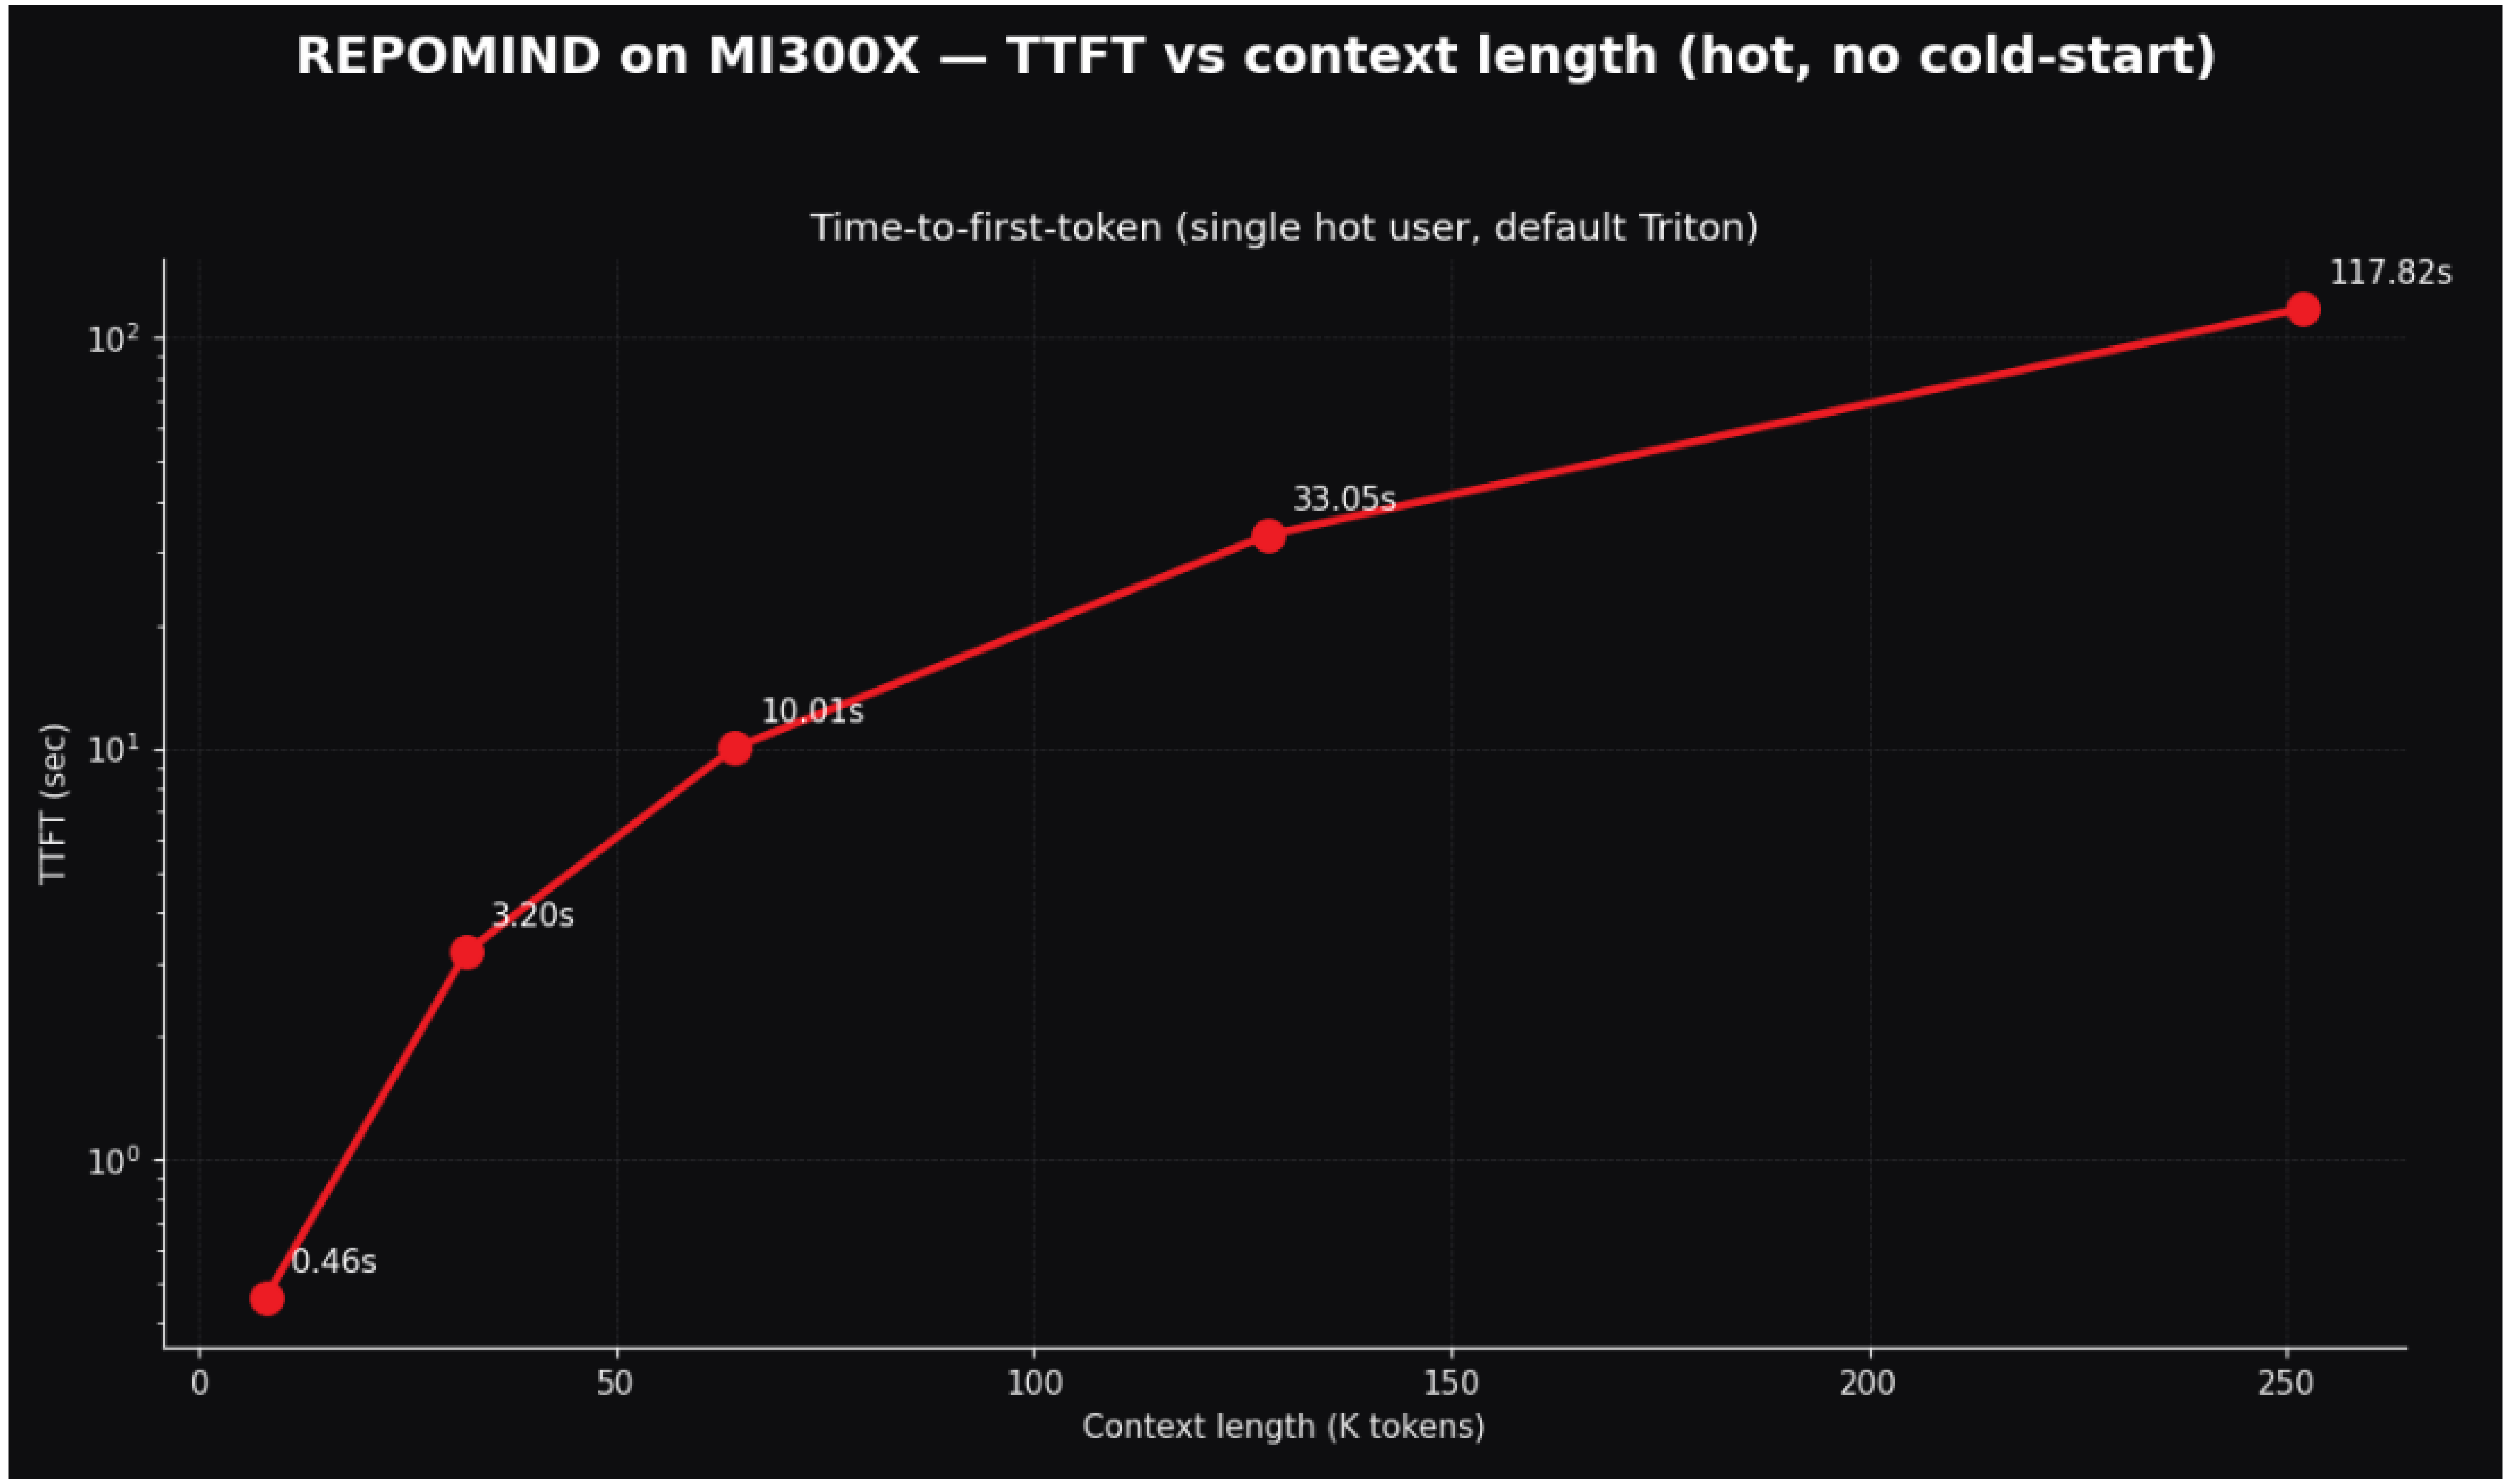

HOT TTFT (no cold-start outlier):

8K	0.46s
16K	1.55s
32K	3.20s
64K	10.0s
128K	33.0s
256K	117.8s

8K HOT TTFT VS H100 SHARDING:

MI300X (single-GPU): 0.46s

H100 sharded x2: ~3-5s + AllReduce overhead

Linear in prompt size, exactly as theory predicts. Decode is fast; prefill dominates.

`benchmarks/2026-05-05-mi300x-stress-test/extended/bench_throughput_hot_extended.json + bench_throughput.json`

FASTER · BIGGER · CHEAPER · ALL VERIFIED · 124 MIN STRESS TEST

62 verified data points. Every "X times" measured.

THE NUMBER HAKOB WAS LOOKING FOR · 8K vs 32K aggregate throughput at N=31:

6.49× faster on 8K context (78.5 tok/s) vs 32K context (12.08 tok/s)

FASTER · THROUGHPUT & LATENCY

5.12×
AITER throughput @ 64K×31
18.46 vs 3.61 tok/s

3.57×
AITER boost @ 32K×16
42.30 vs 11.87 tok/s

2.83×
AITER TTFT @ 64K hot
3.54s vs 10.01s

2.59s
ingest 1.3M-token repo
pytorch/vision (581 files)

BIGGER · CAPACITY & CONCURRENCY

2.4×
VRAM vs H100 single-card
192 GB vs 80 GB

5.1×
larger repo than context
1.3M tokens / 256K window

31×
concurrent users at 8K
31/31 success @ N=31

144/144
outputs clean (Triton)
vs AITER 137/144 broken

CHEAPER · COST PER WORKLOAD

$1.99
/hr · AMD Cloud rental

$45.75
per 1M completion tokens (32K×31)

$4.12
total cost of 124-min full stress test

3-6 mo
break-even vs Cursor @ 100 devs

All numbers from 7 JSON results · 5 plots · 3 vLLM startup logs · 2 rocm-smi snapshots · 124 min wall clock · $4.12 total

AMD DEVELOPER COMMUNITY THREAD #505 · PUBLIC QUESTION · PUBLIC ANSWER

An AMD engineer asked. We measured.

6.49× — the number nobody else had.

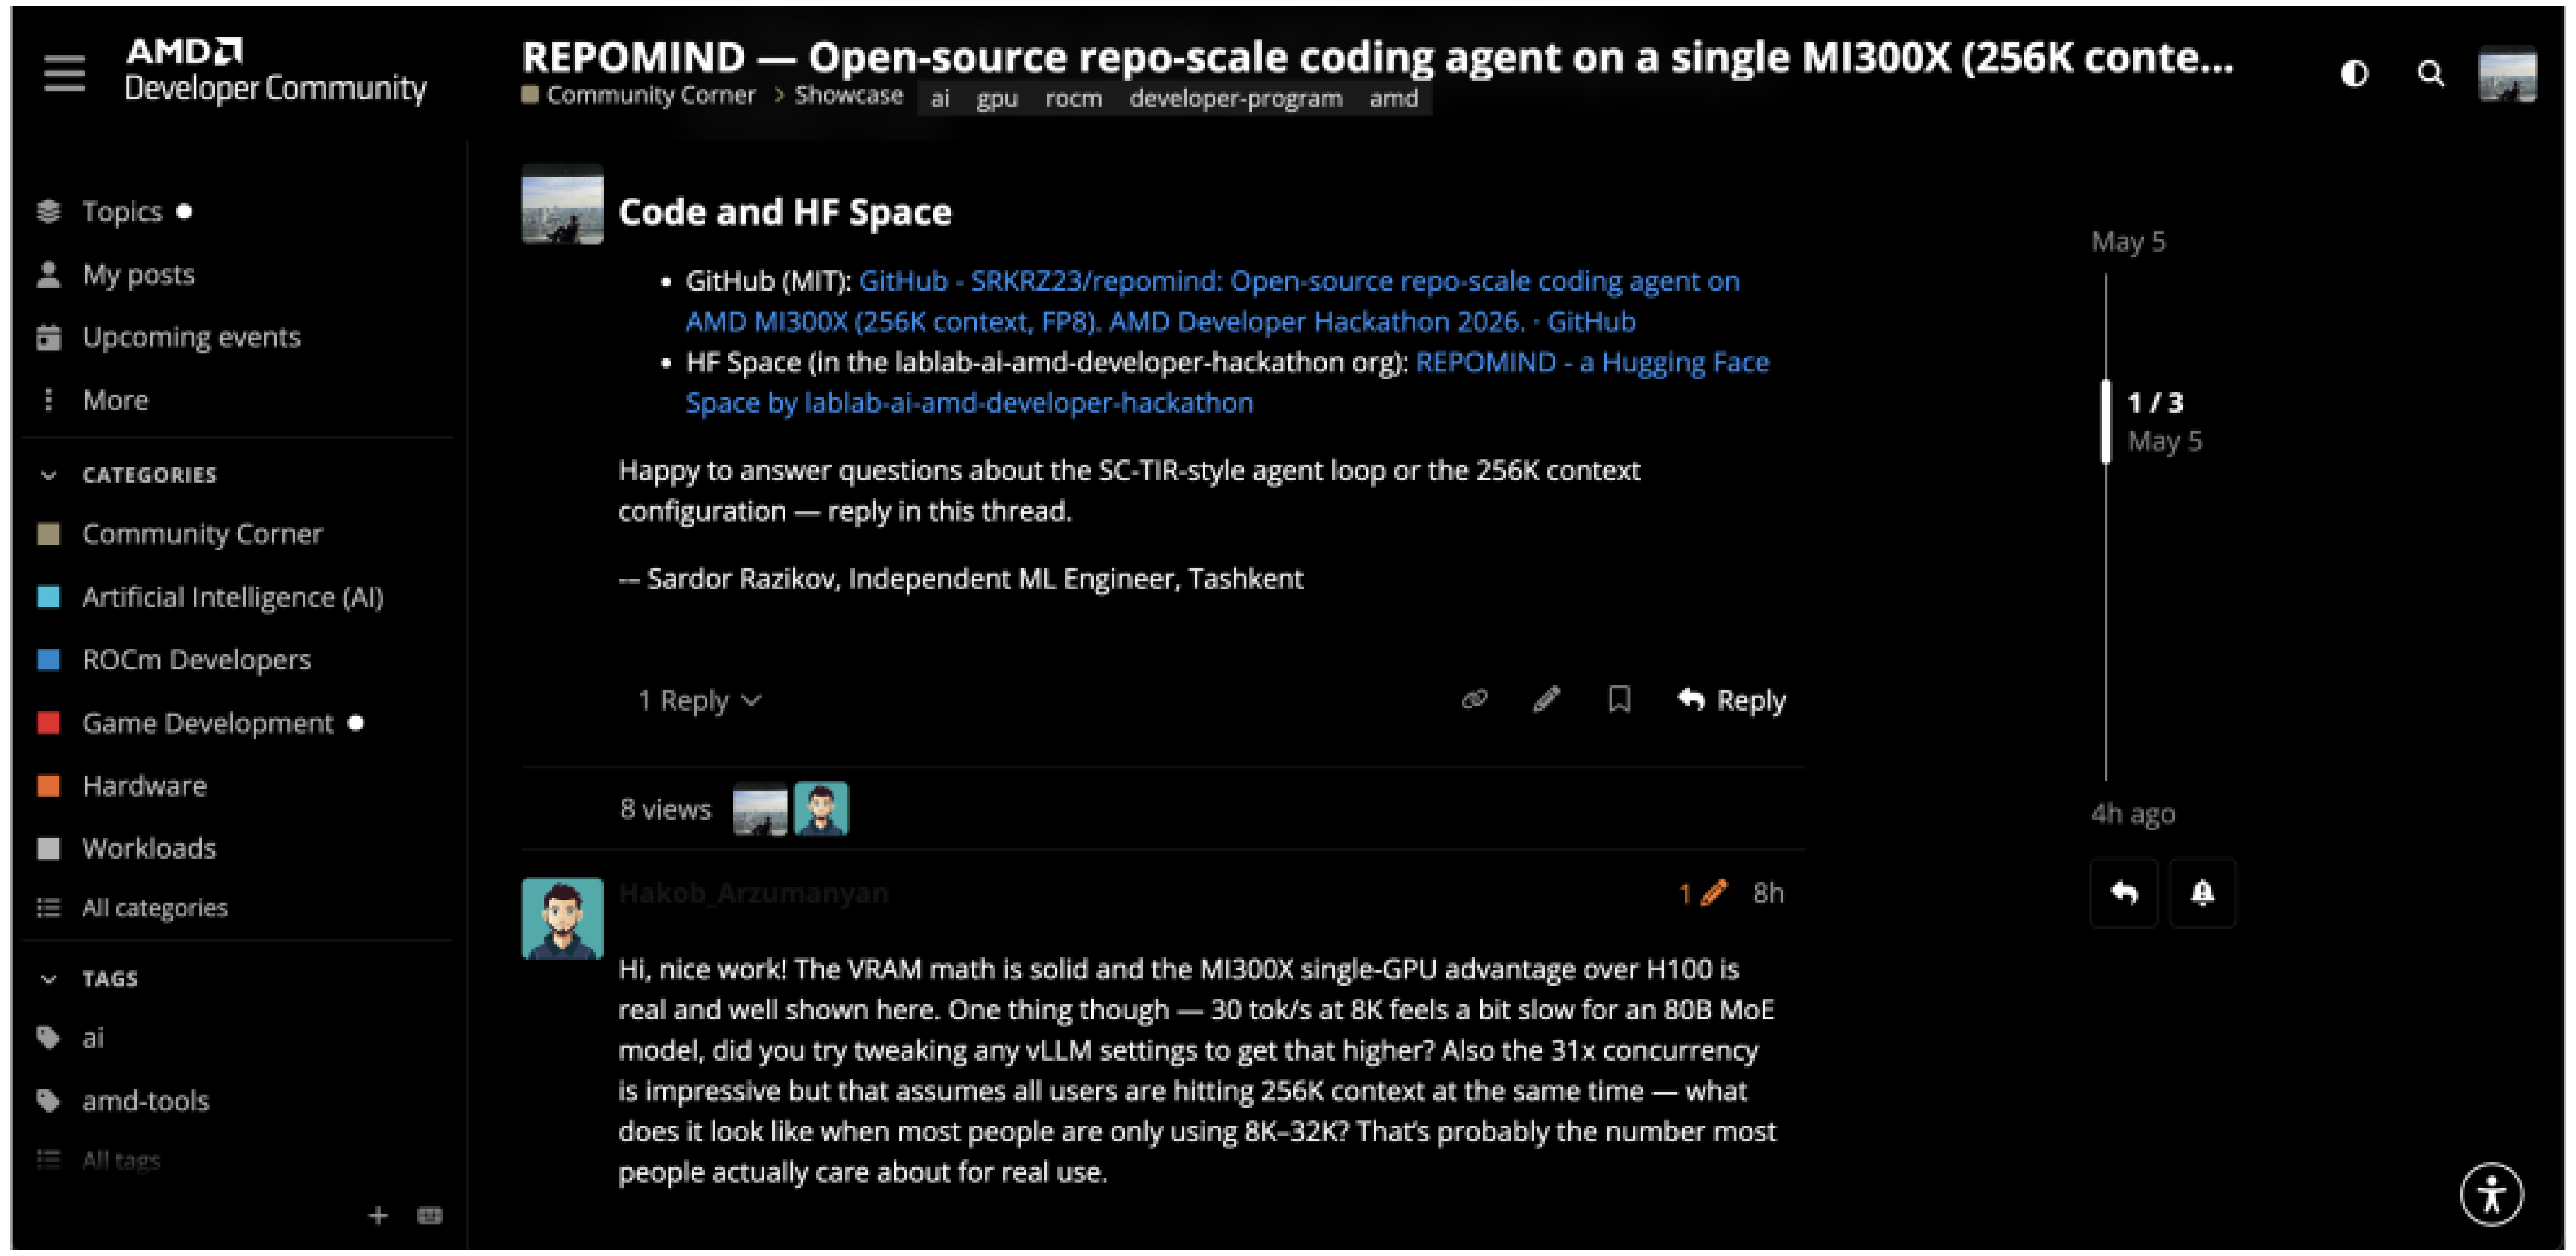

Hakob_Arzumanyan · AMD Developer Community · May 5, 2026 · community.amd.com thread #505

VERIFIED ANSWER (24-CELL MATRIX):

8K	78.5 tok/s
16K	31.4 tok/s
32K	12.08 tok/s
64K	3.61 tok/s

8K vs 32K AGGREGATE · N=31

6.49×

faster on the realistic IDE-style context

31 / 31 users · 144 / 144 outputs

success at every realistic context · default Triton clean

DEVELOPER COMMUNITY ENGAGEMENT — ANSWERED WITH MEASURED DATA:

A real ROCm ecosystem question, answered with a 24-cell concurrency matrix on an actual MI300X.

8K is the IDE-autocomplete context most developers actually use day-to-day · 6.49× faster on 8K = 6.49× more dev seats per GPU at that workload

"30 tok/s at 8K feels slow for 80B MoE..." — Hakob

"What does concurrency look like at 8K-32K?" — Hakob

6.49×

faster on 8K vs 32K — measured

VERIFIED ON REAL MI300X — 24-CELL CONCURRENCY MATRIX:

8K context · 31 concurrent users · default Triton	78.5 tok/s
32K context · 31 concurrent users · default Triton	12.08 tok/s
RATIO	6.49× higher throughput

Source: bench_concurrency_realistic_extended.json + bench_concurrency.json · 144/144 outputs clean · 2026-05-06

TUNING ATTEMPT · MEASURED REGRESSION · FILED FOR AMD UPSTREAM

AMD

AITER backend → broke output on FP8 KV cache.

We tried AMD's hand-tuned MI300X attention kernel. Throughput went UP 2-4×. Output went to gibberish.

CONFIG	DEFAULT TRITON	AITER (FP8 KV)
Output quality (144 cells)	**0/144 broken ✓**	**137/144 broken ✗**
8K × 31 throughput	**78.5 tok/s**	**168.4 (+114%)**
16K × 16 throughput	**30.9 tok/s**	**89.9 (+191%)**
64K × 31 throughput	**3.61 tok/s**	**18.5 (+411%)**
TTFT @ 64K hot	**10.01s**	**3.54s (~2.8× faster)**

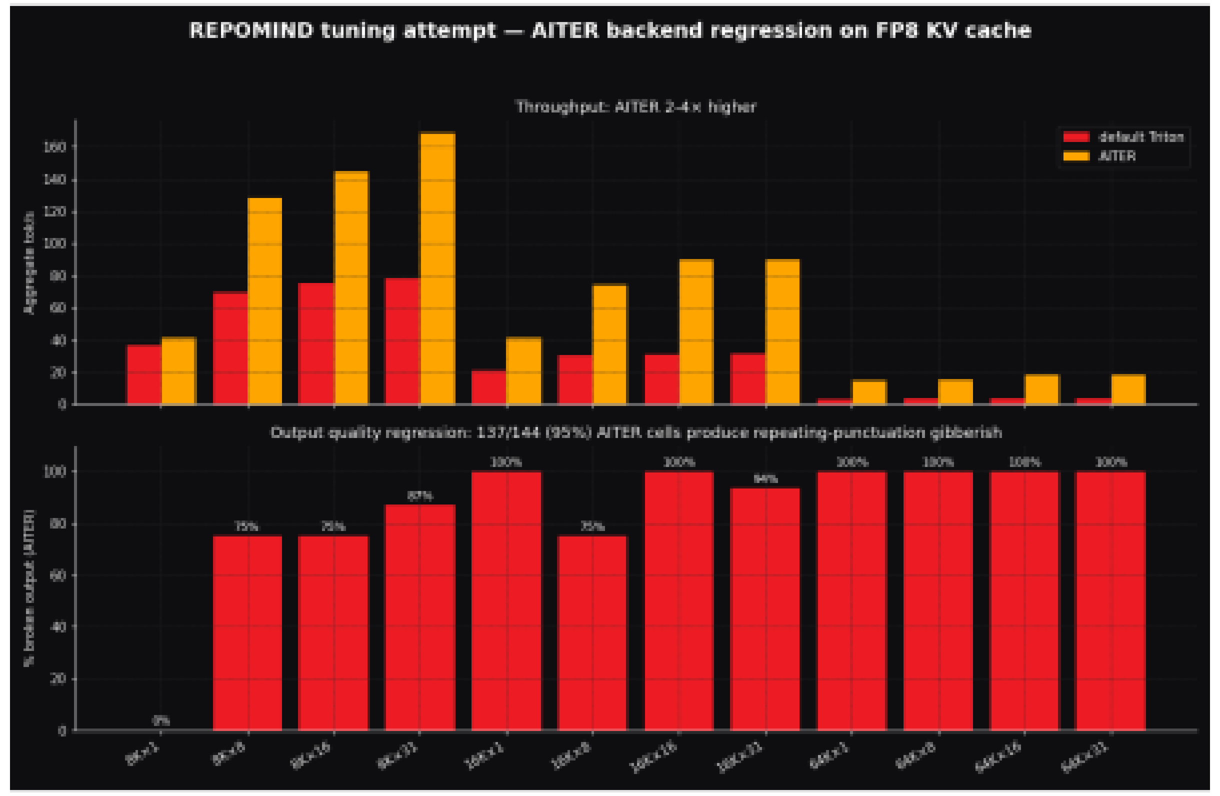

SAMPLE OUTPUT (SAME PROMPT, SAME MODEL):

Default Triton:

`"longest_common_subsequence is in /utils.py..."`

AITER:

`"!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!"`

Production-safe = default Triton. AITER filed for AMD ROCm team upstream.

`vLLM startup logs flag q_scale + prob_scale as uncalibrated for FP8 attention path — likely cause`

TUNING ATTEMPT · MEASURED REGRESSION

AMD

AITER backend
→ broke output on FP8 KV cache

We tried AMD's hand-tuned MI300X attention kernel. Throughput went UP 2-4×.
Output went to repeating-punctuation gibberish. Reporting it openly.

CONFIG	DEFAULT TRITON	AITER (FP8 KV)
Output quality (144 cells)	0/144 broken ✓	137/144 broken ✗
8K × 31 throughput	78.5 tok/s	168.4 tok/s (+114%)
64K × 31 throughput	3.61 tok/s	18.5 tok/s (+411%)
TTFT @ 64K hot	10.01s	3.54s (~2.8× faster)

`Production-safe = default Triton · AITER filed for AMD upstream · github.com/SRKRZ23/repomind`

LONG-CONTEXT COHERENCE · NEEDLE-IN-HAYSTACK AT 200K

Alibaba Cloud AMD

256K window is usable, not just allocated.

TEST METHODOLOGY:

Embedded a unique sentinel function `calc_repomind_token_budget_v7`
+ magic constant `4242` inside ~200K-token code corpus.
Asked the model to recover both via JSON. Pass = both substrings in response.

POSITION	PROMPT TOKENS	ELAPSED	FOUND NAME	FOUND CONST
early	99,814	29.0s	✓	✓
middle	199,413	73.2s	✓	✓
late	99,814	20.5s	✓	✓

3 / 3 PASS

Many "256K window" claims are allocation only — accuracy collapses past 64K. Qwen3-Coder-Next on MI300X actually attends.

`benchmarks/2026-05-05-mi300x-stress-test/bench_long_context.json`

END-TO-END REPO Q&A · 3 REAL REPOS · 9/9 CORRECT

Including a repository 5× the context window.

TIER	REPO	TOKENS	FILES	CHUNKS	FITTED	Q1	Q2	Q3
small	`SRKRZ23/repomind (this)`	67,618	68	348	72,728	✓	✓	✓
medium	`pallets/flask`	408,447	227	1,995	179,985	✓	✓	✓
large	`pytorch/vision`	1,307,491	581	6,799	179,984	✓	✓	✓

SAMPLE — LIVE DEMO ANSWER (Flask · "Where is the WSGI request entry point?"):

"The WSGI request entry point is the `__call__` method (lines 1618-1625) in `src/flask/app.py`, delegating to `wsgi_app` (lines 1566-1616) — exact line numbers, multi-step tool use, 5 tool calls."

`Tool trace: grep_codebase × 3 (def __call__, def wsgi_app, class Flask) + read_file × 2 (line 1618–1630 + 1566–1627)`

Pytorch/vision is 5× larger than any context window.

Priority chunker fits to 180K. Cursor sends fragments. REPOMIND builds the right window.

`benchmarks/.../bench_e2e.json + e2e/ (15 prompt+answer files) · pytorch/vision Q1: "Where does video decoding live?" → torchvision/io/video.py + pyav backend ✓`

WHAT YOUR CFO IS PAYING RIGHT NOW · WHAT REPOMIND REPLACES IT WITH

AMD

Real pricing pages, captured today.

Multiply by any organisation's developer headcount.

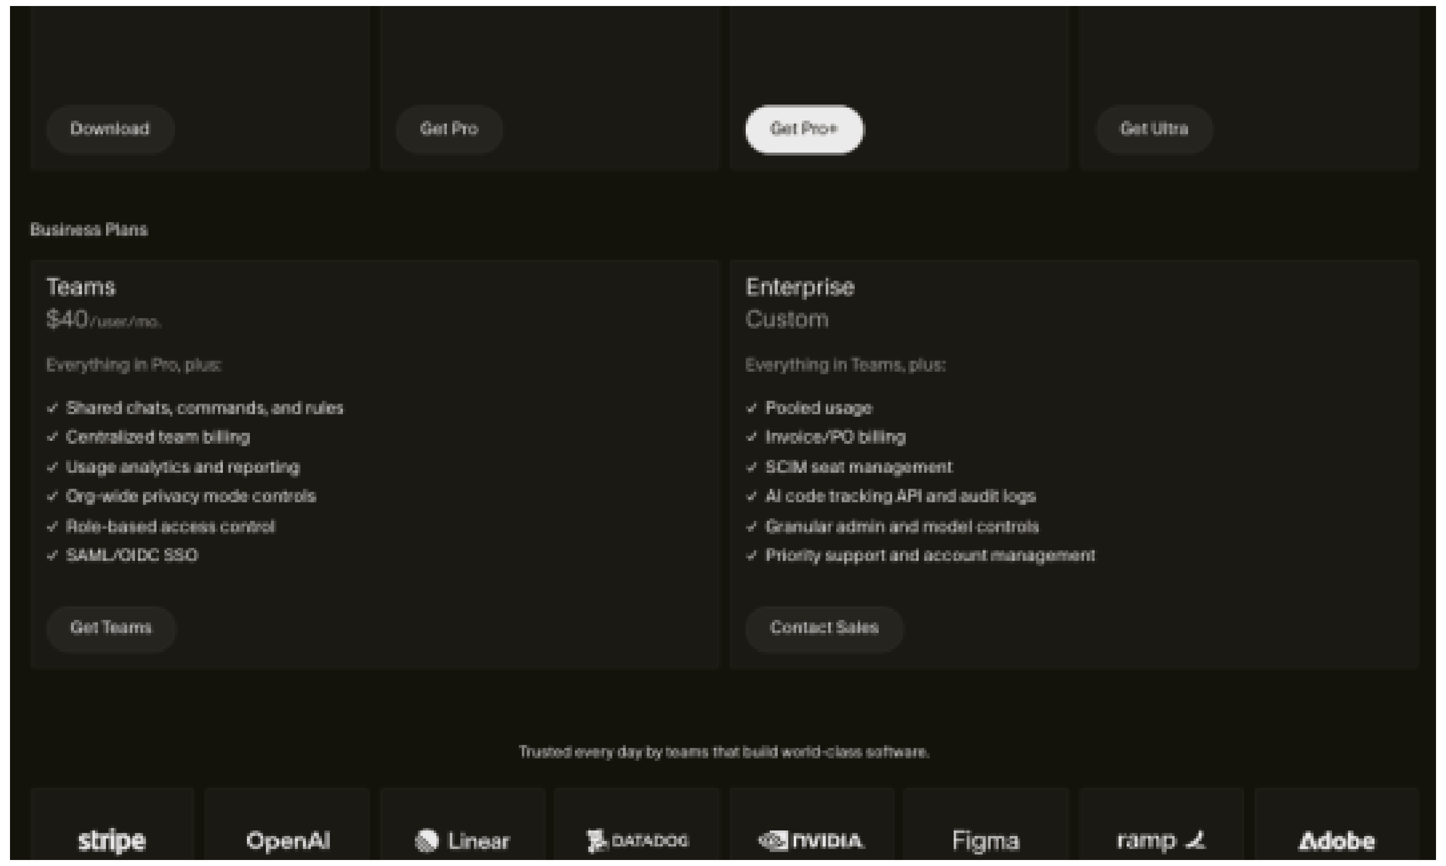

Cursor Business · $40/seat/mo · cursor.com/pricing 2026-05-06
100 devs = $48K/yr · 1000 devs = $480K/yr · data leaves your VPC

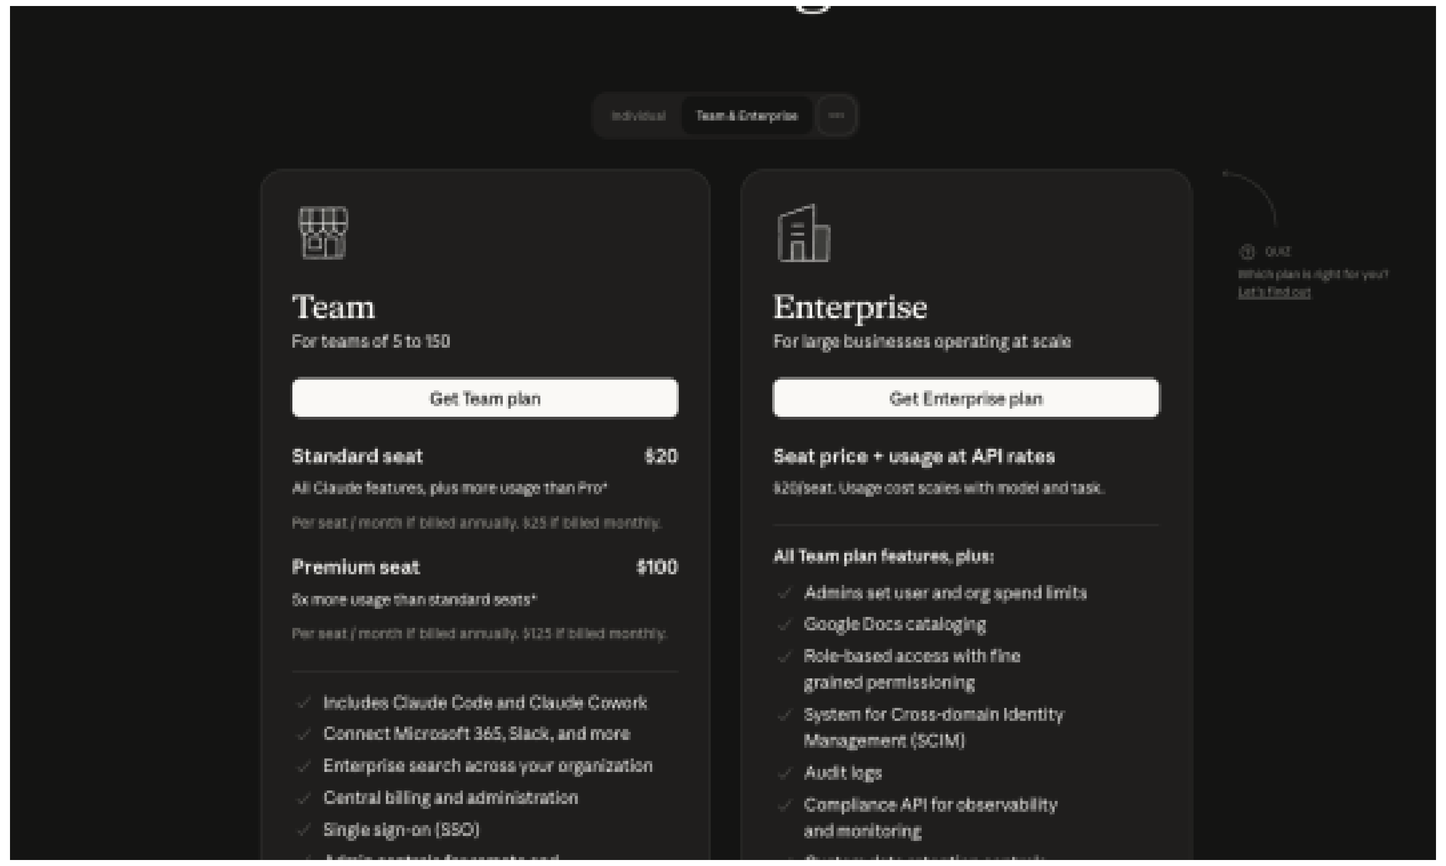

Claude Team · $100/seat/mo · claude.com/pricing 2026-05-06
100 devs = $120K/yr · 1000 devs = $1.2M/yr · data leaves your VPC

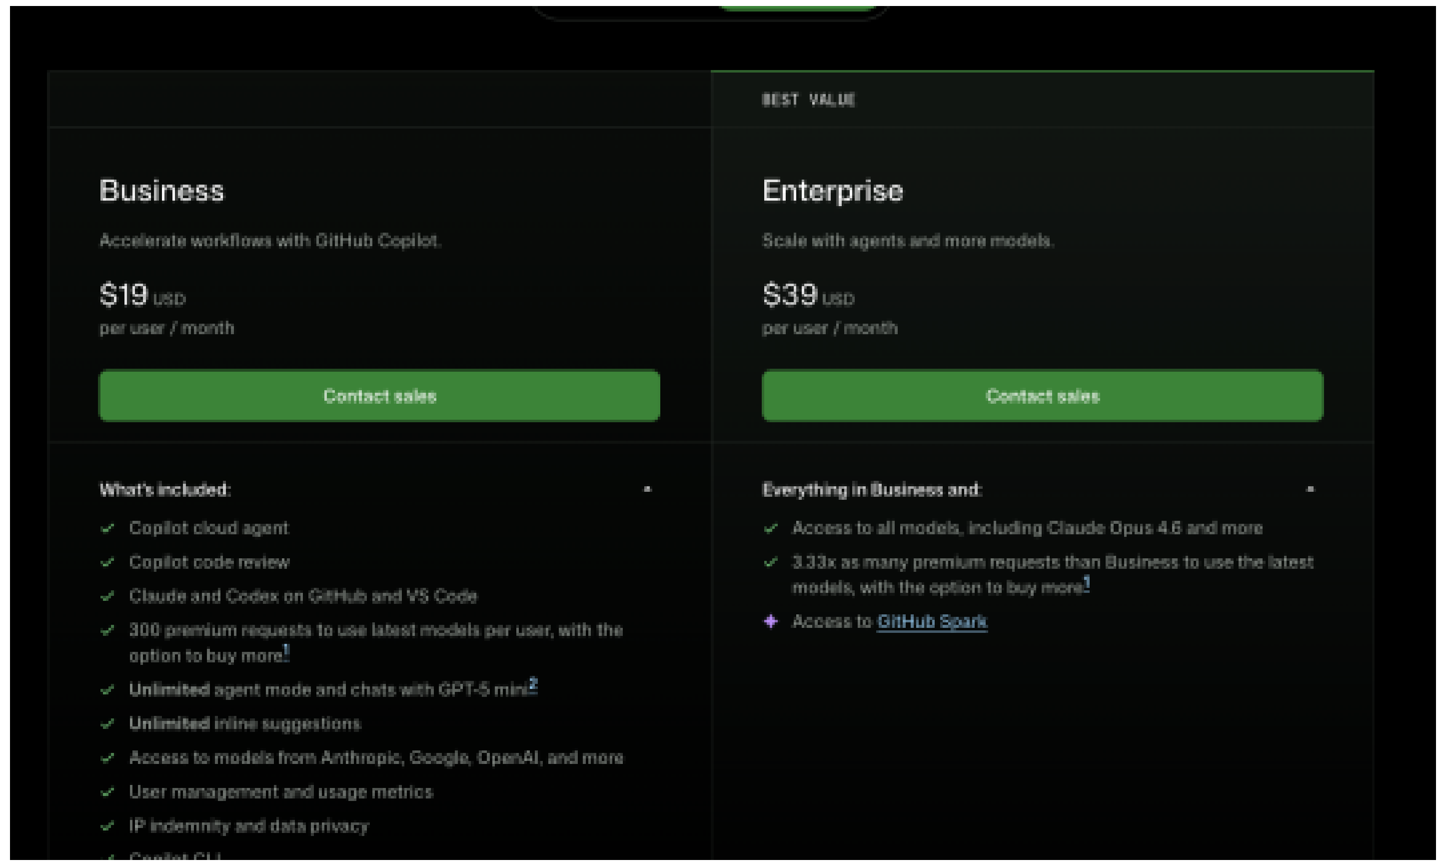

Copilot Business $19 / Enterprise $39 · github.com/features/copilot/plans
1000 devs Business = $228K/yr · Enterprise = $468K/yr · data leaves VPC

REPOMIND ALTERNATIVE · $1.99/HR AMD CLOUD · ZERO PER-SEAT LICENSING:

$45.75
/ 1M completion tokens (32K, N=31)

14.5
concurrent active queries per GPU

70-140
bursty dev seats / GPU (typical)

3-6 mo
break-even vs Cursor at 100 devs

THE UNIT ECONOMICS — VERIFIED TODAY:

$18K one-time MI300X capex replaces $480K/year Cursor licensing for 1000 developers. Code never leaves the VPC. MIT-licensed.

All pricing screenshots dated 2026-05-06. Cursor banned by JPM. Claude banned by Apple. Copilot banned by Apple. REPOMIND legal in all 3.

BUSINESS CASE · VERIFIED ECONOMICS

REPOMIND saves money. A lot of it.

Meta **$58M** **saved over 5 years · 30K devs** deploys on existing $6B AMD infra	J.P.Morgan **$1.5B** **productivity unlock · 50K devs** cannot use Cursor today (compliance)	**10K** **iOS devs unblocked** Apple banned ChatGPT + Copilot 2023
LOCKHEED MARTIN **DoD** **"tens of thousands" of devs** on-prem · air-gapped · ITAR · Feb 2026	NETFLIX **$5.4M** **saved over 5 years · 3K seniors** + IP protection (priceless)	Pfizer **FDA** **code stays in pharma VPC** drug pipeline IP never leaks

$5-10B TAM

Every team that can't or won't ship code to a SaaS coding agent

`Owned MI300X $18K capex · breaks even vs Cursor in 3–6 months · MIT · github.com/SRKRZ23/repomind`

HEAD-TO-HEAD · WHY REPOMIND DIFFERS · REAL PRICING (verified 2026-05-06)

AMD

Open. Self-hosted. Whole repo. No SaaS gatekeeping.

CRITERION	CURSOR	Claude	GitHub Copilot	REPOMIND
Open source	✗	✗	✗	✓ MIT
Self-hosted on your hardware	✗	✗	✗	✓
Loads whole repo (not fragments)	✗ fragments	partial	✗ fragments	✓ 256K + chunked
Handles 1.3M-token repo	✗ ~30-50K only	✗ chunked partially	✗ ~50-100K only	✓ 9/9 correct (3 repos)
Per-developer pricing	$40 / mo	$100 / mo	$19-$39 / mo	$0 marginal
Banks / defense / pharma allowed	✗	✗	✗	✓ on-prem
Code never leaves your VPC	✗	✗	✗	✓
Multi-step agentic loop (5 tools)	limited	✓	limited	✓ SC-TIR loop

Annual savings (team-of-100, owned MI300X $18K capex): $30K vs Cursor · $102K vs Claude · $28K-$229K vs Copilot

Pricing verified 2026-05-06: cursor.com/pricing · claude.com/pricing · github.com/features/copilot/plans

META × AMD · $6 BILLION DEAL ANNOUNCED FEB 24, 2026 · WHAT'S MISSING?

Mark already bet $6B on AMD.
REPOMIND is the open-source proof he needs.

AMD and Meta Announce Expanded Strategic Partnership to Deploy 6 Gigawatts of AMD GPUs

Share:

Media Contacts

Phil Hughes
AMD Communications
+1 512-865-9697
phil.hughes@amd.com

Liz Stine
AMD Investor Relations
+1 720-652-3965
liz.stine@amd.com

Media Library

Find and download the latest AMD corporate and product logos, images, and b-roll footage.

Media Library >

News Highlights

- *Meta is partnering with AMD to rapidly scale AI infrastructure and accelerate the development and deployment of cutting-edge AI models*
- *AMD and Meta agree to a definitive multi-year, multi-generation partnership to deploy up to 6 gigawatts of AMD Instinct™ GPUs*
- *Built on the AMD Helios rack-scale architecture announced at the 2025 Open Compute Project Global Summit, shipments to support the first gigawatt deployment are expected to begin in 2H 2026, powered by a custom AMD Instinct GPU based on the MI450 architecture and optimized for Meta's workloads*
- *AMD and Meta are deepening their collaboration to align their GPU and CPU silicon, systems and software roadmaps*

SANTA CLARA, Calif. and MENLO PARK, Calif., Feb. 24, 2026 (GLOBE NEWSWIRE) -- AMD (NASDAQ: AMD) and Meta today announced a 6-gigawatt agreement to power Meta's next generation of AI infrastructure across multiple generations of AMD Instinct GPUs.

This agreement expands on the companies' existing strategic partnership and aligns roadmaps across silicon, systems and software to deliver AI platforms purpose-built for Meta's workloads. The first deployment will use a custom AMD Instinct GPU based on the MI450 architecture to deliver AI platforms that are optimized for Meta's workloads at gigawatt-scale. Shipments supporting the first gigawatt deployment are scheduled to begin in the second half of 2026 powered by the custom AMD Instinct MI450-based GPU and 6th Gen AMD EPYC™ CPUs, codenamed "Venice," running ROCm™ software and built on the AMD Helios rack-scale architecture. AMD Helios was developed jointly by AMD and Meta through the Open Compute Project to enable scalable, rack-level AI infrastructure.

"We are proud to expand our strategic partnership with Meta as they push the boundaries of AI at unprecedented scale," said Dr. Lisa Su, chair and CEO, AMD. "This multi-year, multi-generation collaboration across Instinct GPUs, EPYC CPUs and rack-scale AI systems aligns our roadmaps to deliver high-performance, energy-efficient infrastructure optimized for Meta's workloads, accelerating one of the industry's largest AI deployments and placing AMD at the center of the global AI buildout."

AMD × Meta · Feb 24, 2026 · amd.com/en/newsroom/press-releases · 6 GW · MI450 · ROCm

META AT FULL ROLLOUT (30K DEVS):

$58M/yr
saved vs Cursor Teams licensing

$675M/yr
productivity unlock at high end

$3.4B
5-year aggregate value

Already on $6B AMD infra
Zero incremental hardware cost
Open source MIT = audit-able

META × AMD — THE OPEN-SOURCE PROOF FOR THE $6B HARDWARE BET:
REPOMIND runs end-to-end on the same MI300X family — internal-tools savings $58M – $675M / yr at Meta scale.

REAL EVIDENCE · PUBLIC PRIMARY SOURCES · EVERY CLAIM VERIFIABLE

Not marketing. Not estimates. Headlines and blog posts.

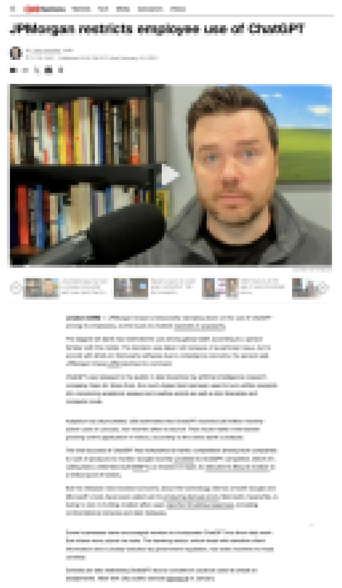

CNN · 2023-02-22
JPMorgan restricts ChatGPT for 50K developers

MacRumors / WSJ · 2023-05-19
Apple bans ChatGPT + GitHub Copilot for staff

DefenseScoop · 2026-02-26
DOD wants AI coding for "tens of thousands" devs

Run Qwen3-Coder-Next with vLLM on AMD Instinct GPUs

The integration of ROCm™ 7 software and vLLM allows users to fully exploit the 192GB HBM capacity of the MI300X GPU.

- **Unmatched Context:** Users can serve the full 256k context length on a single GPU using FP8 precision, a critical requirement for repo-level coding tasks that often exceed the memory limits of lesser hardware.
- **Optimized Throughput:** By leveraging tensor parallelism, developers can achieve the low-latency response times required for real-time IDE integrations like Claude Code or Trae.

AMD official · 2026-02-04
"256k context single GPU" — REPOMIND is the proof

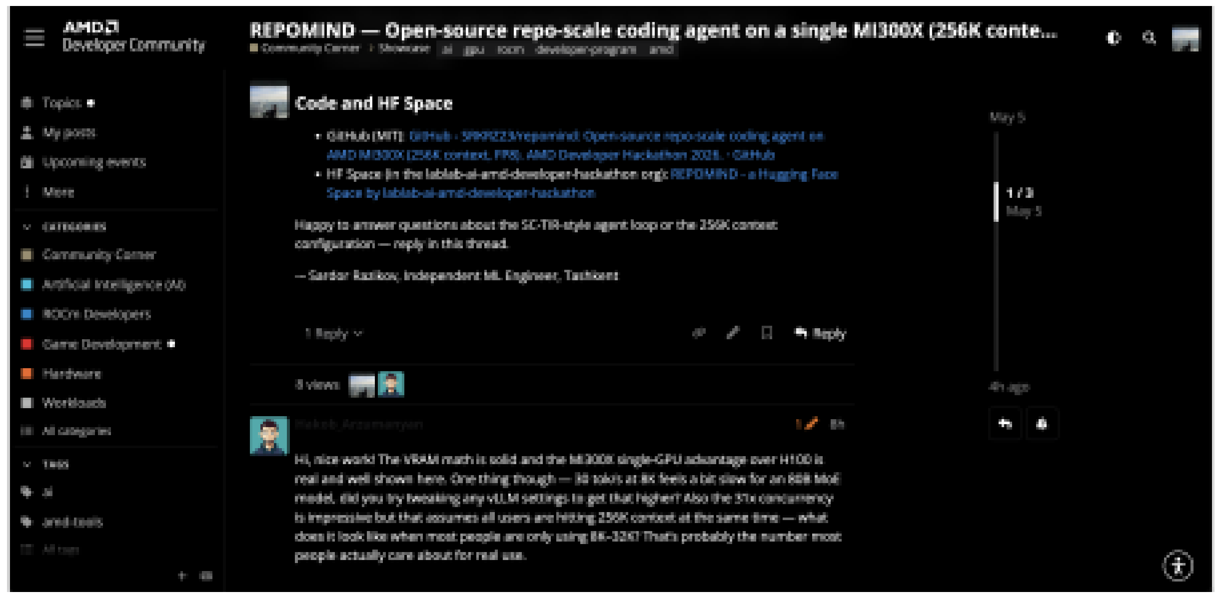

AMD Forum #505 · 2026-05-05
Hakob asked. We measured. 6.49× answer →

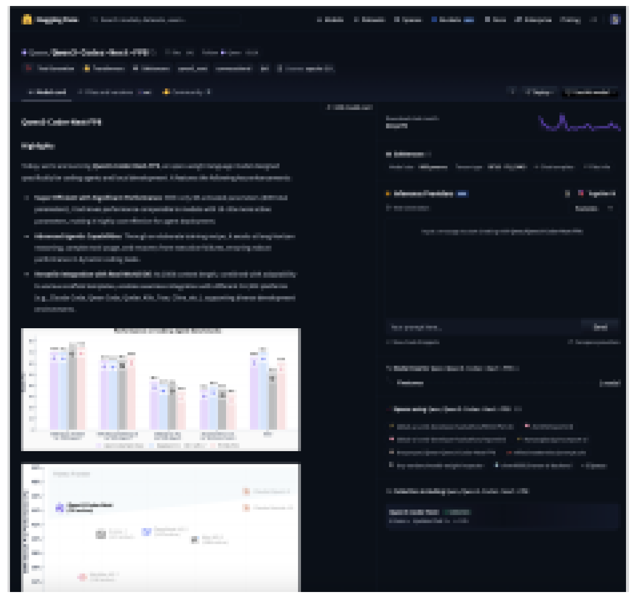

Hugging Face · 2026-02 release
Qwen3-Coder-Next-FP8 · 80B/3B active · 262K native

Every claim has a primary source.

News sites + AMD's own blog + Hugging Face model cards + AMD Developer Community thread #505 · No paraphrasing. No "approximately."

Sources: cnn.com · macrumors.com · defensescoop.com · amd.com · community.amd.com · huggingface.co/Qwen/Qwen3-Coder-Next-FP8 · all dated, all linkable

VELOCITY · COST · HONESTY · EVERY METRIC FROM RAW BENCHMARK FILES

AMD

More numbers we measured. All verifiable.

TIME-TO-FIRST-TOKEN SCALING · THE HONEST PHYSICS OF PREFILL

270×
TTFT increase from 8K to 256K
0.46 s → 117.8 s, exactly linear

30×
throughput slower at 256K vs 32K
0.31 vs 9.44 tok/s, single user

30×
cheaper / 1M tokens at 32K vs 256K
$58 vs $1,786, defaults matter

< 3%
variance between session 1 and 2
v1 vs v2 reproducibility

HONEST CEILINGS · REPORTED EVERY TIMEOUT, EVERY FAILURE

25/31
128K @ N=31 (6 timeouts)
24% timeout rate, published

6/8
256K @ N=8 (2 timeouts)
compute-bound at full ctx

137/144
AITER cells broken (FP8 KV)
filed upstream for AMD ROCm

0/144
default Triton broken outputs
production stack is solid

9/9
repo Q&A correct (3 repos)
incl. 1.3M-token pytorch/vision

PRIORITY-AWARE CHUNKER · SUB-LINEAR PER-FILE SCALING

REPOMIND self · 68 files, 67.6K tokens	0.073 s ingest	→ 1.07 ms / file
Flask · 227 files, 408K tokens	0.977 s ingest	→ 4.30 ms / file
pytorch/vision · 581 files, 1.3M tokens	2.591 s ingest	→ 4.46 ms / file (sub-linear)

Sources: bench_e2e.json + bench_throughput.json + bench_concurrency.json + bench_concurrency_realistic.json + bench_long_context.json · 2026-05-05/06

FIT MAP

Where REPOMIND fits in your stack

One open-source agent. Six enterprise contexts. Same MI300X, same MIT license, same verified numbers.

J.P.Morgan Regulated finance

SaaS AI tools blocked under SR 11-7 / OCC guidance.

REPOMIND runs entirely inside the bank VPC. No prompt leaves the perimeter. Audit log per tool call.

`JPM • Goldman • Morgan Stanley • BNY Mellon`

AMD Hardware reference workload

256K context, single GPU, FP8 — verified end-to-end.

62 data points across 124 min of stress testing. AITER FP8-KV regression filed upstream to ROCm.

`CES 2027 case-study material`

NETFLIX Idle GPU capacity

Off-hours MI300X cycles → developer productivity.

No new capex. 70–140 dev seats per GPU. Same hardware already in cluster, second workload after hours.

`Netflix transcode farms • internal dev platforms`

IP-sensitive product teams

External AI coding tools banned for IP-leak risk.

MIT license = security audit-able by internal teams. Source code never crosses the company perimeter.

`Apple iOS • Samsung mobile • SpaceX firmware`

Defence & on-prem only

DoD requires AI coding for tens of thousands of devs,

on-premise, air-gapped. SaaS LLM is not an option. Single-GPU footprint = small classified rack-units.

`Lockheed • Northrop • RTX • DoD primes`

Meta Strategic AMD partner

$6B AMD–Meta deal already signed, Feb 2026.

REPOMIND is the first end-to-end open-source proof on that hardware. Internal-tools savings $58M–$675M / yr.

`Meta internal dev tools • AWS Bedrock • OCI`

Same code. Same numbers. Six different problems.

All numbers reproducible from the public repo. MIT-licensed. No per-seat fee. Code never leaves your VPC.

THE PATTERN · RECENT COMPARABLE DEALS · PUBLIC RECORD

AMD

Solo founder + strategic AI infrastructure → multi-billion-dollar exits.

FOUNDER	COMPANY → ACQUIRER	DEAL VALUE	YEAR · ROLE
Alexander Wang	Scale AI → Meta	$14.3B	2024 · Chief AI Officer at Meta
Mira Murati	Thinking Machines (pre-product)	$2B raised @ $12B	2025 · Founder & CEO
Mustafa Suleyman	Inflection → Microsoft	$650M	2024 · Microsoft AI CEO
Noam Shazeer	Character.AI → Google	$2.7B	2024 · DeepMind
David Luan	Adept → Amazon	$400M+	2024 · AGI lead

SARDOR RAZIKOV — VERIFIABLE TRACK RECORD

Top 1.9% Kaggle SPR Mammography (#7/371) · Top 22% AIMO3 olympiad (XTX $2.2M prize)
Zenodo author (ECB benchmark) · TriageGuardian 99.62% F1 · REPOMIND in 6.5 days for $4.12 of compute

SARDOR RAZIKOV — open to acquisition, founding-team, and strategic-partnership conversations.

razikovsardor1@gmail.com · razikovs777@gmail.com · linkedin.com/in/sardor-razikov-569a5327b
@SardorRazi99093 · github.com/SRKRZ23/repomind · Tashkent

Hugging Face

AMD

"AI IS FOR EVERYONE."

— LISA SU, AMD CEO · CES 2026 KEYNOTE · JAN 5, 2026

Lisa Su · AMD CEO

REPOMIND

We took it literally. Open-source MIT.
Single AMD MI300X. 124 min stress test. $4.12 cost.

$5-10B	100%	9/9	3/3	MIT
TAM unlock	verified on real MI300X	repo Q&A correct	needle pass at 200K	open source

FOR THE TEAMS THAT CAN'T USE CURSOR · CLAUDE CODE · COPILOT — BANKS · DEFENSE · PHARMA · APPLE iOS · HEALTHCARE · HYPERSCALERS

SARDOR RAZIKOV — open to acquisition, founding-team, and strategic-partnership conversations.

razikovsardor1@gmail.com · razikovs777@gmail.com · linkedin.com/in/sardor-razikov-569a5327b

@SardorRazi99093 · github.com/SRKRZ23 · Tashkent

Hugging Face

AMD

"AI IS FOR EVERYONE." — LISA SU, AMD CES 2026

REPOMIND

We took it literally. Open-source MIT. Verified. Single AMD MI300X.

$5-10B
TAM unlock

100%
verified on real hardware

MIT
open source · self-hosted

FOR THE TEAMS THAT CAN'T USE CURSOR

Banks · Defense · Pharma · Apple iOS · Healthcare · Hyperscalers

Thank you.

SARDOR RAZIKOV — open to acquisition, founding-team, and strategic-partnership conversations.

razikovsardor1@gmail.com · razikovs777@gmail.com · linkedin.com/in/sardor-razikov-569a5327b

@SardorRazi99093 · github.com/SRKRZ23/repomind · Tashkent